

1 The digital image

- Problems of digital cameras sensors
- transmission interference
 - compression artefacts
 - spilling
 - scratches, sensor noise
 - bad contrast

1.1 Image as 2D signal

Signal: function depending on some variable with physical meaning
Image: continuous function
2 variables: xy -coordinates
3 variables: $xy+time$


Brightness is usually the value of the function, but other physical values are: Temperature, pressure, depth...

- What is an image?**
- A picture or pattern of a value varying in space and/or time
 - Representation of a function $f: \mathbb{R}^n \rightarrow S$
 - In digital form, e.g.:
 $I: \{1, \dots, X\} \times \{1, \dots, Y\} \rightarrow S$.
 - For greyscale CCD images, $n = 2, S = \mathbb{R}^+$

What is a pixel?
Not a little square! E.g. analog or cubic reconstruction filter.

1.2 Image sources

digital camera (CCD)



analog to digital Conversion

- The ADC measures the charge and digitizes the result.
 - Conversion happens line by line.
 - The charges in each photosite move down through the sensor array.
- Blooming** Buckets have finite capacity. Photosite saturation causes blooming.
- Bleeding or smearing** during transit buckets still accumulate some charges. Due to tunneling and CCD Architecture. (Influenced by time "in transit" versus integration time. Effect is worse for short shutter times)

dark current CCDs produce thermally-generated charge. They give non-zero output even in darkness. Partly, this is the *dark current* and it fluctuates randomly. One can reduce it by cooling the CCD.

CMOS Has same sensor elements as CCD. Each photo sensor has its *own amplifier*. This leads to more noise (reduced by subtracting "black" image) and lower sensitivity (lower fill rate). The uses of standard CMOS technology allows to put other components on chip and "smart" pixels.

CCD vs. CMOS *CCD:* mature technology, specific technology, high production cost, high power consumption, higher fill rate, blooming, sequential readout.
CMOS: recent technology, standard IC technology, cheap, low power, less sensitive, per pixel amplification, random pixel access, smart pixels, on chip integration with other components, rolling shutter (sequential read-out of lines)

1.3 Sampling

1D Sampling takes a function and returns a vector whose elements are values of that function at the sample points.

Undersampling "Missing" things between samples. Information lost

aliasing signals "traveling in disguise" as other frequencies. (Can happen in undersampling.)

1.4 Reconstruction

Inverse of sampling. Making samples back into continuous function. For output (need reliable method), for analysis or processing (need mathematical method), amounts to "guessing" what the function did in between.

Bilinear interpolation

$$f(x,y) = (1-a)(1-b)f[i,j] + a(1-b)f[i+1,j] + abf[i+1,j+1] + (1-a)bf[i,j+1]$$

Nyquist frequency Half the sampling frequency of a discrete signal processing system. Signal's max frequency (bandwidth) must be *smaller* than this.

sampling grids cartesian sampling, hexagonal sampling and non-uniform sampling

1.5 Quantization

real valued function will get digital values - integer values. Quantization is lossy and can't be reconstructed. Simple quantization uses equally spaced levels with k intervals.

usual quantization intervals

Grayscale image: 8 bit = $2^8 = 256$ gray-values. *Color image RGB (3 channels):* 8bit/channel = $2^{24} = 16.7M$ colors.

1.6 Image Properties

Image resolution: Clipped when reduced. **Geometric resolution:** Whole picture but crappy when reduced. **Radio-metric resolution:** Number of colors.

1.7 Image Noise

additive Gaussian Noise Common model $I(x,y) = f(x,y) + c$, where $c \sim \mathcal{N}(0, \sigma^2)$. So that $p(c) = (2\pi\sigma^2)^{-1} e^{-c^2/2\sigma^2}$.

Poisson noise (shot noise) $p(k) = \lambda^k e^{-\lambda} / k!$

Rician noise (appears in MRI)

$$p(I) = \frac{I}{\sigma^2} \exp\left(-\frac{(I^2 + f^2)}{2\sigma^2}\right) I_0\left(\frac{If}{\sigma^2}\right)$$

Multiplicative noise: $I = f + fc$
Signal to noise ration (SNR) $s = F/\sigma$ is an index of image quality, where

$$F = \frac{1}{XY} \sum_{x=1}^X \sum_{y=1}^Y f(x,y)$$

Often used instead: *Peak Signal to Noise Ratio (PSNR)* $s_{\text{peak}} = F_{\text{max}}/\sigma$

1.8 Colour Images

Consist of red, green and blue channel.

Prism (with 3 sensors) Separate light in three beams using dichroic prism. Requires 3 sensors and precise alignment. Gives good color separation. → high-end cameras

Filter mosaic Coat filter directly on sensor. "Demosaiicing" to obtain full colour & full resolution image. → low-end cameras

Filter wheel rotate multiple filters in front of lens. Allows more than 3 colour bands. → static scenes

new color CMOS sensor, foveon's X3 blue, green, red sensor, one above the other (descending) → better image quality

2 Image segmentation

"Segmentation is the ultimate classification problem. Once solved, Computer Vision is solved."

Interim Summary Segmentation is hard. It is easier if you define the task carefully: (1) Segmentation task binary or continuous? (2) What are regions of interest? (3) How accurately must the algorithm locate the region boundaries?

1.5 Quantization

real valued function will get digital values - integer values. Quantization is lossy and can't be reconstructed. Simple quantization uses equally spaced levels with k intervals.

usual quantization intervals

Grayscale image: 8 bit = $2^8 = 256$ gray-values. *Color image RGB (3 channels):* 8bit/channel = $2^{24} = 16.7M$ colors.

1.6 Image Properties

Image resolution: Clipped when reduced. **Geometric resolution:** Whole picture but crappy when reduced. **Radio-metric resolution:** Number of colors.

threshold T :

$$B(x,y) = \begin{cases} 1 & \text{if } I(x,y) \geq T \\ 0 & \text{if } I(x,y) < T. \end{cases}$$

Choosing T By trial and error. Compare results with ground truth. Automatic methods. (ROC curve)

Chromakeying Control Lighting! "Plain" distance measure (e.g.)

$$I_\alpha = |I - g| > T$$

$$T \sim 20, g = \begin{pmatrix} 0 & 255 & 0 \end{pmatrix}^T$$

Gaussian model per pixel (Like chromakeying) mean $\mu \rightarrow I_\mu$, standard deviation $\sigma \rightarrow I_\Sigma$. $I_\alpha = |I - I_\mu| > T$

$$T = \begin{pmatrix} 20 & 20 & 10 \end{pmatrix}, I_{bg} = \text{background image. Or better (e.g.)}$$

$$I_\alpha = \sqrt{(I - I_{bg})^T \Sigma^{-1} (I - I_{bg})} > T$$

ROC Analysis Receiver operating Characteristic. An ROC curve characterizes the performance of a binary classifier. A binary classifier distinguishes between two different types of things.

Classification error Binary classifiers make errors. Two types of input to a binary classifier: Positives, negatives. Four possible outcomes in any test: True positive, true negative, false negative, false positive.

ROC Curve Characterizes the error trade-off in binary classification tasks. It plots the *true positive fraction*

$$TP \text{ fraction} = \text{true positive count} / P,$$

$$P = TP + FN \text{ and false positive fraction}$$

FP fraction = false positive count/N, $N = FP + TN$. ROC curve always passes through (0,0) and (1,1).

MAP (Maximum A Posteriori) detector

Operating points choose an *operating point* by assigning relative costs and values to each outcome, $V_{TN}, V_{TP}, C_{FN}, C_{FP}$. V and C being values and costs. For simplicity, often $V_{TN} = V_{TP} = 0$.

Performance Assessment In real-life, we use two or even three separate sets of test data: (1) A *training set*, for tuning the algorithm, (2) A *validation set* for tuning the performance score, (3) An *unseen test set* to get a final performance score on the tuned algorithm.

Pixel connectivity Define neighbors, e.g. (for 2D) 4-neighborhood or 8-neighborhood

Pixel paths There are e.g. 4- and 8-connected paths. (p_i neighbor of p_{i+1}).

Connected regions A region is 4- or 8-connected if it contains a(n) 4- or 8-connected path between any two of its pixels.

2.2 Region Growing

(1) Start from a seed point or region. (2) Add neighboring pixels that satisfy the criteria defining a region. (3) Repeat

until we can include no more pixels.

```
function B = RegionGrow(I, seed)
[X,Y] = size(I);
visited = zeros(X,Y);
visited(seed) = 1;
boundary = emptyQ;
while (~boundary.empty())
    nextPoint = boundary.deQ();
    if (include(nextPoint, visited))
        visited(nextPoint) = 1;
        foreach (x,y) in N(nextPoint)
            if (visited(x,y) == 0)
                boundary.enQ(x,y);
                visited(x,y) = 1;
        end
    end
end
```

2.2.1 Variations

seed selection (1) One seed point, (2) Seed region, (3) Multiple seeds.

seed selection (1) Greylevel thresholding, (2) Greylevel distribution model. E.g. include if $(I(x,y) - \mu)^2 < (n\sigma)^2$, $n = 3$. Can update μ and σ after every iteration, (3) color or texture information.

snakes A snake is an *active contour*. It's a polygon. Each point on contour moves away from seed while its image neighborhood satisfies an inclusion criterion. Often the contour has smoothness constraints, the algorithm iteratively minimizes an energy function:

$$E = E_{\text{tension}} + E_{\text{stiffness}} + E_{\text{image}}$$

2.3 Spatial relations

Markov Random Fields Markov chains have 1D structure. At every time, there is one state. This enabled use of dynamic programming. *Markov Random fields* break this 1D structure: • Field of sites, each of which has a label, simultaneously. • Label at one site dependend on others, no 1D structure to dependencies. • This means no optimal, efficient algorithms, except for 2-label problems. Minimize

$$\text{Energy}(y; \text{data}) = \sum_i \psi_1(y_i; \text{data}_i) + \sum_{i,j} \psi_2(y_i, y_j; \text{data}_i, \text{data}_j)$$

• $i, j \in \text{edges}$

FG-BG segmentation The code does the following: • Background RGB Gaussian model training (from many images) • shadow modeling (hard shadow and soft shadow) • graphcut foreground-background segmentation

2.4 Morphological Operations

They are local pixel transformations for processing region shapes. Most often used on binary images. Logical transformations based on comparison of pixel neighborhoods with a pattern.

8-neighbor erode (Minkowsky subtraction) Erase any foreground pixel that has one eight-connected neighbor that is background

8-neighbor dilate (Minkowsky addition) Paint any background pixel that has one eight-connected neighbor that is foreground. *Applications:* Smooth region boundaries for shape analysis, remove noise and artefacts from an imperfect segmentation, match particular pixel configurations in an image for simple object recognition

structuring elements morphological operations take two arguments 1. a binary image 2. a structuring element Compare the structuring element to the neighborhood of each pixel. This determines the output of the morphological operation. The structuring element is also a binary array and has an origin.

$$I_1 \cup I_2 = \{x: x \in I_1 \text{ or } x \in I_2\}, I_2 \cap I_1 = \{x: x \in I_1 \text{ and } x \in I_2\}, I^c = \{x: x \notin I\}, I_1 \setminus I_2 = \{x: x \in I_2 \text{ and } x \notin I_1\}.$$

Erosion of binary image I by the structuring element S is defined by $I \ominus S = \{z \in E \mid S_z \subset I\}$

S_z translation of S by vector z .

Dilation is $I \oplus S = \bigcup_{b \in S} I_b$.

Opening $I \circ S = (I \ominus S) \oplus S$.

Closing $I \bullet S = (I \oplus S) \ominus S$.

To remove holes in the foreground and islands in the background, do both opening and closing. Thesize and shape of the structuring element determine which features survive. In the absence of knowledge about the shape of features to remove, use a circular structuring element.

Granulometry Provides a size distribution of distinct regions or "granules" in the image. We open (opening as above) the image with increasing structuring element size and count the number of regions after each operation. Creates "morphological sieve".

```
function gSpec = granulo(I, T, maxRad)
% Segment the image I.
B = (I > T);
%Open the image at each structuring element size up
%to a maximum and count the remaining regions.
for x=1:maxRad
    O = imopen(B, strel('disk', x));
    numRegions(x) = max(max(connectedComponents(O)));
end
gSpec = diff(numRegions);
```

Sequential thinning/thickening With structuring elements S_1, \dots, S_n and sequential thinning/thickening $I \blacklozenge \{S_i: i = 1, \dots, n\} = ((I \blacklozenge S_1) \cdots \blacklozenge S_n)$

Several sequences of structuring elements are useful in practice. These are usually the set of rotations of a single structuring element, sometimes called the *Golay alphabet*. See *bwmorph* in *matlab*

2.4.1 Medial Axis Transform (MAT, skeletonization)

The skeleton and MAT are stick-figure representations of a region $X \in \mathbb{R}^2$. Start a grassfire at the boundary of the region, theskeleton is the set of points at which two fire fronts meet.

Skeleton Use structuring element

$$B = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{pmatrix}.$$

The n -th skeleton subset is $S_n(X) = (X \ominus nB) \setminus [(X \ominus nB) \odot B]$, where $\ominus n$ denotes n successive erosions. The skeleton is the union of all the skeleton subsets $S(X) = \bigcup_{n=1}^{\infty} S_n(X)$.

Reconstruction can reconstruct region X from its *skeleton subsets*.

$$X = \bigcup_{n=1}^{\infty} S_n(X) \odot nB$$

Applications and problems The skeleton/MAT provies a stick figure representing the region shape. Used in object recognition, in particular, character recognition. *Problems:* Definition of a maximal disc is poorly defined on a digital grid and is sensitive to noise on the boundary. Sequential thinning output sometimes preferred to skeleton/MAT.

3 Image filtering

Image filtering is modifying the pixels in an image based on some function of a local neighborhood of the pixels.

3.1 Linear Shift-Invariant Filtering

About modifying pixels based on *neighborhood*. Local methods simplest. Linear means *linear combination* of neighbors. Linear methods simplest. *Shift-invariant* means doing the same for each pixel. Same for all is simplest. Useful to: Low-level image processing operations, smoothing and noise reduction, sharpen, detect or enhance features.

Linear operation L is a *linear* operation if

$$L[\alpha I_1 + \beta I_2] = \alpha L[I_1] + \beta L[I_2]$$

Output I' of linear image operation is a weighted sum of each pixel in the input I

$$I'_j = \sum_{i=1}^N \alpha_i I_{ij}, \quad j = 1 \cdots N$$

Linear Filtering Linear operations can be written:

$I'(x,y) = \sum_{i,j \in N(x,y)} K(x,y;i,j) I(i,j)$
 $I =$ input image, k is *kernel* of the operation. $N(m,n)$ is a neighbourhood of (m,n) .

Correlation e.g. template matching.

Linear operation: $I' = KI$

$$I'(x,y) = \sum_{i,j \in N(x,y)} K(i,j) I(x+i,y+j)$$

Convolution e.g. point spread function
 $I'(x,y) = \sum_{i,j \in N(x,y)} K(i,j) I(x-i,y-j)$

Edge The filter window falls off the edge of the image, we need to extrapolate, methods: (1) clip filter (black) (2) wrap around (3) copy edge (4) reflect across edge (5) vary filter near edge

Filter at boundary (1) ignore, copy or truncate. No processing of boundary pixels. Pad image with zeros (matlab). Pad image with copies of edge rows/columns (2) truncate kernel (3) reflected indexing (4) circular indexing

Separable Kernels Separable filters can be written

$$K(m,n) = f(m)g(n)$$

for a rectangular neighborhood with size $(2M+1) \times (2N+1)$,

$$I'(m,n) = f * (g * I(N(m,n))),$$

$$I''(m,n) = \sum_{j=-N}^N g(j) I(m,n-j),$$

$$I''(m,n) = \sum_{i=-M}^M f(i) I''(m-i,n).$$

→ $(2M+1) + (2N+1)$ operations!

Smoothing kernels (low-pass filters)

• Mean filter: $\frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$, • Weighted

$$\text{smoothing filters: } \frac{1}{10} \begin{pmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{pmatrix}.$$

Gaussian Kernel Idea: Weight contributions of neighboring pixels:

$$\mathcal{N}_{\mu=0, \sigma}(x,y) = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Smoothing with a Gaussian instead of a box filter removes the artefact of the vertical and horizontal lines. Gaussian smoothing Kernel is *separable*! $\mathcal{N}(x,y) = \mathcal{N}(x) \mathcal{N}(y)$. Amount of smoothing depends on σ and window size. Width $> 3\sigma$.

Scale space Convolution of a Gaussian with σ with itself is a gaussian with $\sigma\sqrt{2}$. Repeated convolution by a Gaussian filter produces the scale space of an image.

Gaussian filter top-5 (1) Rotationally symmetric. (2) Has a single lobe. → Neighbor's influence decreases monotonically. (3) Still one lobe in frequency domain. → No corruption from high frequencies (4) Simple relationship to σ (5) Easy to implement efficiently

Differential filters • Prewitt operator:

$$\begin{pmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{pmatrix}, \quad \bullet \text{ Sobel operator: } \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix}.$$

High-pass filters • Laplacian operator: $\begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix}$, • High-pass filter: $\begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}.$

Differentiation and convolution

df/dx = lim_{epsilon -> 0} ((f(x+epsilon,y) - f(x,y))/epsilon)

df/dx approx (f(xn+1,y) - f(xn,y))/Delta x

which is obviously a convolution (-1 1)

Filters and templates Filters at some point can be seen as taking a dot-product between the image and some vector, the image is a set of dot products, filters look like the effects they are intended to find, filters find effects they look like.

Image sharpening Also known as enhancement. Increases the high frequency components to enhance edges.

I' = I + alpha |K * I|, where K is a high-pass filter kernel and alpha in [0, 1].

Integral images integral images (also known as summed-area tables) allow to efficiently compute the convolution with a constant rectangle

S(x,y) = integral_0^x integral_0^y I(x',y')

A = S(1) A+C = S(3)

A+B = S(2) A+B+C+D = S(4)

D = S(4) - S(2) - S(3) + S(1)

Also possible along diagonal.

Viola-Jones cascade face detection

Very efficient face detection using integral images.

4 Image features

4.1 Template matching

Problem Locate an object, described by a template t(x,y), in the image s(x,y). Example: Passport photo as image and eyes to detect.

Method Search for the best match by minimizing mean squared error E(p,q)

= sum_{x,y=-inf}^inf [s(x,y) - t(x-p,y-q)]^2 = sum_{x,y=-inf}^inf |s(x,y)|^2 + |t(x,y)|^2 - 2 * sum_{x,y=-inf}^inf s(x,y) * t(x-p,y-q)


Equivalently, maximize area correlation

r(p,q) = sum_{x,y=-inf}^inf s(x,y) * t(x-p,y-q) = s(p,q) * t(-p,-q) <= sqrt([sum [s(x,y)]^2] * [sum [t(x,y)]^2])

where in the last step the Cauchy-Schwarz inequality was used. Equality <=>

s(x,y) = alpha * t(x-p,y-q) with alpha >= 0 Area correlation is equivalent to convolution of image s(x,y) with impulse response t(-x,-y).

Diagram of template matcher



4.2 Edge detection

Idea (continuous-space): Detect local gradient

||nabla(f(x,y))|| = sqrt((df/dx)^2 + (df/dy)^2)

Digital image: Use finite differences instead:

difference (-1 1), central difference (-1 0 1); Prewitt

Sobel

Roberts

Laplacian operator Detects discontinuities by considering second derivative

Nabla^2 f(x,y) = d^2 f/dx^2 + d^2 f/dy^2

Isotropic (rotationally invariant) operator, zero-crossings mark edge location, discrete-space approximation by convolution with 3 x 3 impulse response

Laplacian of Gaussian The Laplacian operator is very sensitive to fine detail and noise, so blur it first with Gaussian. -> do it in one operator

Laplacian of Gaussian (LoG)

LoG(x,y) = -1/(pi * sigma^4) * [1 - (x^2 + y^2)/(2 * sigma^2)] * e^(-x^2 + y^2/(2 * sigma^2))

Find points for which the following is min Delta^T M Delta

for ||Delta|| = 1. i.e. maximize eigenvalues of M.

Keypoint detection Often based on eigenvalues lambda_1, lambda_2 of M ("structure matrix" = normal matrix) / second-moment matrix")

M(x,y) = sqrt((df/dx)^2 + (df/dy)^2)

alpha(x,y) = arctan (df/dy / df/dx)

(3) Apply nonmaxima suppression to gradient magnitude image (4) Double thresholding to detect strong and weak edge pixels (5) Reject weak edge pixels not connected with strong edge pixels

Canny nonmaxima suppression Quantize edge normal to one of four directions: horizontal, -45°, vertical, 45°. If M(x,y) is smaller than either of its neighbors in edge normal direction -> suppress; else keep

Double-thresh. of grad. magn.

strong edge: M(x,y) >= theta_high

weak edge: theta_high > M(x,y) >= theta_low

Typical setting: theta_high, theta_low = 2,3. Region labeling of edge pixels. Reject regions without strong edge pixels.

4.3 Feature detection

4.3.1 Hough transform

Problem: fit a straight line (or curve) to a set of edge pixels. Hough transform (1962): generalized template matching technique. (1) Consider detection of straight lines y = mx + c. (2) draw a line in the parameter space m,c for each edge pixel x,y and increment bin counts along line. Detect peak(s) in (m,c) plane. (3) Alternative parametrization avoids infinite-slope problem x cos theta + y sin theta = rho

circle detection find circles of fixed radius r. For circles of undetermined radius.

dius, use 3d Hough transform for parameters (x0,y0,r)

4.3.2 Detecting corner points

Many applications benefit from features localized in (x,y). Edges well localized only in one direction -> detect corners. Desirable properties of corner detector: (1) Accurate localization, (2) invariance against shift, rotation, scale, brightness change, (3) robust against noise, high repeatability

4.3.3 Most accurately localizable patterns

Local displacement sensitivity

S(Delta x, Delta y) = sum_{x,y in window} [f(x,y) - f(x - Delta x, y + Delta y)]

Linear approximation for small Delta x, Delta y

f(x + Delta x, y + Delta y) approx f(x,y) + df/dx (x,y) Delta x + df/dy (x,y) Delta y

S(Delta x, Delta y) approx sum_{(x,y) in window} [(df/dx df/dy) (Delta x) (Delta y)]

= (Delta x Delta y) M (Delta x Delta y)

SSD approx Delta^T M Delta

Find points for which the following is min Delta^T M Delta

for ||Delta|| = 1. i.e. maximize eigenvalues of M.

Keypoint detection Often based on eigenvalues lambda_1, lambda_2 of M ("structure matrix" = normal matrix) / second-moment matrix")

M = sum_{(x,y) in window} [(df/dx df/dy) (df/dx df/dy)]

C(x,y) = det(M) - k * (trace M)^2 = lambda_1 lambda_2 - k * (lambda_1 + lambda_2)

Measure of "cornerness"

C(x,y) = det(M) - k * (trace M)^2 = lambda_1 lambda_2 - k * (lambda_1 + lambda_2)

Represent function on a new basis. Basis elements have the form e^{-i2pi (ux+vy)}

The Fourier transform is

f-hat(u,v) = integral integral_{R^2} dx dy f(x,y) e^{-i2pi (ux+vy)}

Basis functions of Fourier transform are eigenfunctions of linear systems.

Important functions

spatial domain frequency domain

"time" domain "frequency" domain

Impulse, or "delta" function

comb

Boisac

Sync. Func

Compute subpixel localization by fitting parabola to cornerness function

M(x,y) = sum_{x,y in window} G(x-x0,y-y0,sigma)

Corner importance weight Give more importance to central pixels by using Gaussian weighting function

M(x,y) = sum_{x,y in window} G(x-x0,y-y0,sigma)

Compute subpixel localization by fitting parabola to cornerness function

Robustness of Harris corner detector (1) Invariant to brightness offset:

f(x,y) -> f(x,y) + c (2) Invariant to shift and rotation (3) Not invariant to scaling

4.3.4 Lowe's SIFT features

Recover features with position, orientation and scale.

Position (1) Look for strong responses of DoG filter, (2) only consider local maxima.

Scale (1) Look for strong responses of DoG filter over scale space. (2) only consider local maxima in both position and scale. (3) Fit quadratic around maxima for subpixel accuracy.

Orientation (1) Create histogram of local gradient directions computed at selected scale. (2) Assign canonical orientation at peak of smoothed histogram. (3) Each key specifies stable 2D coordinates (x,y,scale,orientation)

SIFT description (1) Thresholded image gradients are sampled over 16 x 16 array of locations in scale space. (2) Create array of orientation histograms (3) 8 orientations x 4 x 4 histogram array = 128 dimensions

5 Fourier Transform

5.1 Aliasing

One can't shrink an image by taking every second pixel. If we do, characteristic errors appear. Typically, small phenomena look bigger; fast phenomena can look slower. Common phenomena (1) Wagon wheels rolling the wrong way in movies. (2) Checkboards misrepresented in ray tracing (3) Striped shirts look funny on color television.

5.2 Definition

Represent function on a new basis. Basis elements have the form e^{-i2pi (ux+vy)}

The Fourier transform is

f-hat(u,v) = integral integral_{R^2} dx dy f(x,y) e^{-i2pi (ux+vy)}

Basis functions of Fourier transform are eigenfunctions of linear systems.

Important functions

spatial domain frequency domain

"time" domain "frequency" domain

Impulse, or "delta" function

comb

Boisac

Sync. Func

Compute subpixel localization by fitting parabola to cornerness function

M(x,y) = sum_{x,y in window} G(x-x0,y-y0,sigma)

Corner importance weight Give more importance to central pixels by using Gaussian weighting function

M(x,y) = sum_{x,y in window} G(x-x0,y-y0,sigma)

Compute subpixel localization by fitting parabola to cornerness function

Robustness of Harris corner detector (1) Invariant to brightness offset:

f-hat = f-hat * f-hat

The Fourier transform of the product of two functions is the convolution of the Fourier transforms

f-hat * g-hat = F-hat(f * g)

5.3 Sampling


Go from continuous world to discrete world, from function to vector. Samples are typically measured on regular grid. We want to be able to approximate integrals sensibly -> Delta function

S_{2D}(f(x,y)) = sum_{i,j=-inf}^inf f(x,y) delta(x-i,y-j)

= f(x,y) sum_{i,j=-inf}^inf delta(x-i,y-j)

with S = Sample operator.

FT of sampled signal



Problem: F-hat[h](u) = 1/h(u). sinc has many zeroes and these frequencies can't be recovered! Solution: Regularized reconstruction filter

F-hat[h](u,v) = F-hat[h] / ||F-hat||^2 + epsilon

Singularities are avoided by the rugelazization epsilon.

Space-time super-resolution One can put two movies of the same thing and merge their frames for space and time super-resolution.

Proper sampling To avoid this effect, this is the procedure: original signal -> lp filtering -> lp filt. sign

sample -> sampl.sign. -> reconstr. -> reconstr.sign

Smoothing as low-pass filtering The message of the FT is that high frequencies lead to trouble with sampling. Solutions press high frequencies before sampling. A filter whose FT is a box is bad, because the filter kernel has infinite support. Common solution: use a Gaussian.

Nyquist sampling theorem Nyquist theorem: The sampling frequency must be at least twice the highest frequency. omega_s >= 2*omega_h. If this is no the case, the signal needs to be bandlimited before sampling, e.g. with a low-pass filter.

5.4 Image Restoration

Pixelization Possibilities: Square pixels, Gaussian reconstruction filter, Bilinear interpolation, perfect reconstruction filter.

Motion blurring Each light dot is transformed into a short line along the x1-axis:

h(x1,x2) = 1/(2*delta) [theta(x1+delta) - theta(x1-delta)] delta(x2)

Noise Gaussian blurring kernel:

h(x1,x2) = 1/(2*pi*sigma^2) exp(-(x1^2 + x2^2)/(2*sigma^2))

General approach (1) Sort samples f(x,y) of an M x N image (or rectangular block in the image) into column vector of length M x N. (2) Compute transform coefficients c. Af where A is a matrix of size (MN)^2. (3) Transform A is unitary, iff A^{-1} = A^* (4) If A is real-valued, i.e. A = A-bar, transform is orthonormal.

pensate the effect of the image degradation h(x), i.e.,

(h * h)(x) = delta(x)

h-hat(x) may be determined more easily in Fourier space:

F-hat[h-hat](u,v) * F-hat[h](u,v) = 1

To determine F-hat[h-hat], we need to estimate (1) the distortion model h(x) (point spread function) or F-hat[h](u,v) (modulation transfer function) (2) the parameters of h(x), e.g. for defocussing

Motion Blur FT

F-hat[h](u,v) = 1/(2*delta) integral_{-delta}^delta dx1 exp(-i2*pi*u*x1)

= sinc(2*pi*u*delta)

Problem: F-hat[h](u) = 1/h(u). sinc has many zeroes and these frequencies can't be recovered! Solution: Regularized reconstruction filter

F-hat[h](u,v) = F-hat[h] / ||F-hat||^2 + epsilon

Singularities are avoided by the rugelazization epsilon.

Space-time super-resolution One can put two movies of the same thing and merge their frames for space and time super-resolution.

Proper sampling To avoid this effect, this is the procedure: original signal -> lp filtering -> lp filt. sign

sample -> sampl.sign. -> reconstr. -> reconstr.sign

Smoothing as low-pass filtering The message of the FT is that high frequencies lead to trouble with sampling. Solutions press high frequencies before sampling. A filter whose FT is a box is bad, because the filter kernel has infinite support. Common solution: use a Gaussian.

Nyquist sampling theorem Nyquist theorem: The sampling frequency must be at least twice the highest frequency. omega_s >= 2*omega_h. If this is no the case, the signal needs to be bandlimited before sampling, e.g. with a low-pass filter.

5.4 Image Restoration

Pixelization Possibilities: Square pixels, Gaussian reconstruction filter, Bilinear interpolation, perfect reconstruction filter.

Motion blurring Each light dot is transformed into a short line along the x1-axis:

h(x1,x2) = 1/(2*delta) [theta(x1+delta) - theta(x1-delta)] delta(x2)

Noise Gaussian blurring kernel:

h(x1,x2) = 1/(2*pi*sigma^2) exp(-(x1^2 + x2^2)/(2*sigma^2))

General approach (1) Sort samples f(x,y) of an M x N image (or rectangular block in the image) into column vector of length M x N. (2) Compute transform coefficients c. Af where A is a matrix of size (MN)^2. (3) Transform A is unitary, iff A^{-1} = A^* (4) If A is real-valued, i.e. A = A-bar, transform is orthonormal.

Energy conservation ||c||^2 = c^* c = F^* A^* Af = ||f||^2

Image collection fi one image, F = f1, ..., fn,

Auto-correlation function R_ff = E[f_i f_i^*] = FF^* / n

energy distribution Energy is conserved, but often will be unevenly distributed among coefficients. Autocorrelation matrix:

R_cc = E[cc^*] = E[Af f^* A^*] = AR_ff A^*

Mean squared values ("average energies") of the coefficients ci are on the diagonal of R_cc.

E[ci^2] = [R_cc] = [AR_ff A^*]_ii

Eigenmatrix of autocorrelation matrix Definition: Eigenmatrix Phi of autocorrelation matrix R_ff. (1) phi is unitary (2) The columns of Phi form a set of eigenvectors of R_ff, i.e.,

R_ff Phi = Phi Lambda

where Lambda = diag(lambda_0, ..., lambda_{MN-1}). (3) R_ff is symmetric nonnegative definite, hence lambda_i >= 0 for all i (4) R_ff is normal matrix, i.e. R_ff R_ff^* = R_ff^* R_ff, hence unitary eigenmatrix exists.

Problem: F-hat[h](u) = 1/h(u). sinc has many zeroes and these frequencies can't be recovered! Solution: Regularized reconstruction filter

F-hat[h](u,v) = F-hat[h] / ||F-hat||^2 + epsilon

Singularities are avoided by the rugelazization epsilon.

Space-time super-resolution One can put two movies of the same thing and merge their frames for space and time super-resolution.

Proper sampling To avoid this effect, this is the procedure: original signal -> lp filtering -> lp filt. sign

sample -> sampl.sign. -> reconstr. -> reconstr.sign

Smoothing as low-pass filtering The message of the FT is that high frequencies lead to trouble with sampling. Solutions press high frequencies before sampling. A filter whose FT is a box is bad, because the filter kernel has infinite support. Common solution: use a Gaussian.

Nyquist sampling theorem Nyquist theorem: The sampling frequency must be at least twice the highest frequency. omega_s >= 2*omega_h. If this is no the case, the signal needs to be bandlimited before sampling, e.g. with a low-pass filter.

5.4 Image Restoration

Pixelization Possibilities: Square pixels, Gaussian reconstruction filter, Bilinear interpolation, perfect reconstruction filter.

Motion blurring Each light dot is transformed into a short line along the x1-axis:

h(x1,x2) = 1/(2*delta) [theta(x1+delta) - theta(x1-delta)] delta(x2)

Noise Gaussian blurring kernel:

h(x1,x2) = 1/(2*pi*sigma^2) exp(-(x1^2 + x2^2)/(2*sigma^2))

General approach (1) Sort samples f(x,y) of an M x N image (or rectangular block in the image) into column vector of length M x N. (2) Compute transform coefficients c. Af where A is a matrix of size (MN)^2. (3) Transform A is unitary, iff A^{-1} = A^* (4) If A is real-valued, i.e. A = A-bar, transform is orthonormal.

Energy conservation ||c||^2 = c^* c = F^* A^* Af = ||f||^2

Image collection fi one image, F = f1, ..., fn,

Auto-correlation function R_ff = E[f_i f_i^*] = FF^* / n

energy distribution Energy is conserved, but often will be unevenly distributed among coefficients. Autocorrelation matrix:

R_cc = E[cc^*] = E[Af f^* A^*] = AR_ff A^*

Mean squared values ("average energies") of the coefficients ci are on the diagonal of R_cc.

E[ci^2] = [R_cc] = [AR_ff A^*]_ii

Eigenmatrix of autocorrelation matrix Definition: Eigenmatrix Phi of autocorrelation matrix R_ff. (1) phi is unitary (2) The columns of Phi form a set of eigenvectors of R_ff, i.e.,

R_ff Phi = Phi Lambda

where Lambda = diag(lambda_0, ..., lambda_{MN-1}). (3) R_ff is symmetric nonnegative definite, hence lambda_i >= 0 for all i (4) R_ff is normal matrix, i.e. R_ff R_ff^* = R_ff^* R_ff, hence unitary eigenmatrix exists.

Problem: F-hat[h](u) = 1/h(u). sinc has many zeroes and these frequencies can't be recovered! Solution: Regularized reconstruction filter

F-hat[h](u,v) = F-hat[h] / ||F-hat||^2 + epsilon

Singularities are avoided by the rugelazization epsilon.

6.2 Basis images and eigenimages (EI)

For a unitary transform, the inverse transform f = A^* c can be interpreted in terms of the superpositions of "basis images" (columns of A^*) of size MN. If the transform is a KL transform, the basis images, which are the eigenvectors of the autocorrelation matrix R_ff, are called "eigenimages". If energy concentration works well, only a limited number of eigenimages is needed to approximate a set of images with small error. These eigenimages form an optimal linear subspace of dimensionality J.

EI for recognition To recognize complex patterns (e.g., faces), large portions of an image (say of size MN) might have to be considered. High dimensionality of "image space" means high computational burden for many recognition techniques. Transform c = Wf can reduce dimensionality from MN to J by representing the image by J coefficients. Idea: tailor a KLT to the specific set of images of the recognition task to preserve the salient features.

Simple recognition Simple Euclidean distance (SSD) between images. Best match wins

argmin Di = ||Ii - I||

Computationally expensive, i.e. requires presented image to be correlated with every image in the database!

Eigenspace matching Let Ii be the input image, I the database. The "character" of the face j = J - (I), with J being any image (set). Do KLT (aka PCA) transformation

Ii -> pi, E^* Ii = pi

Ii - I = Ii - I approx E(pi - p)

Ii - I approx ||pi - p||

with closest rank-k approximation property of SVD. Approximate

argmin Di = ||Ii - I|| approx ||pi - p||

6.3 Eigenfaces (EF)

Concatenate face pixels into "observation vector", x.

EI for recognition (1) Input image, (2) normalize, (3) subtract mean face, (4) KLT, (5) Find most similar pi, (6) similarity measure, (7) rejection system, (8) result of identification.

Limitations of EFs Differences due to varying illumination can be much larger than differences between faces!

6.4 Fisherfaces/LDA

classification task

$$E_{\text{opt}} = \underset{E}{\operatorname{argmax}}(\det(ERE^*)),$$

Fisher linear discriminant analysis: Maximize between-class scatter, while minimizing within-class scatter,

$$F_{\text{opt}} = \underset{F}{\operatorname{argmax}} \left(\frac{\det(FR_BW^*)}{\det(FR_WF^*)} \right),$$
$$R_B = \sum_{i=1}^c N_i (\mu_i - \mu)(\mu_i - \mu)^*,$$
$$R_W = \sum_{\substack{l=1, \dots, c \\ \Gamma_l \in \text{Class}(l)}} (\Gamma_l - \mu_l)(\Gamma_l - \mu_l)^*.$$

N_i are the samples in class i and μ_i is the mean in class i . Solution: Generalized eigenvectors w_i corresponding to the k largest eigenvalues $\{\lambda_i \mid i = 1, \dots, k\}$, i.e.

$R_B w_i = \lambda_i R_W w_i, i = 1, \dots, k$

. Problem: within-class scatter matrix R_W at most of rank $L - c$, hence usually singular. Apply KLT first to reduce dimension of feature space to $L - c$ (or less), proceed with Fisher LDA in low-dimensional space.

Eigenfaces vs. Fisherfaces Eigenfaces conserve energy but the two classes e.g. in 2D are no longer distinguishable. FLD (Fisher LDA) separates the classes by choosing a better 1D subspace. Fisher faces are much better in varying illuminations.

Varying illumination (FF) All images of same Lambertian surface with different illumination (without shadows) lie in a 3D linear subspace. Single point source at infinity

$$f(x, y) = a(x, y) \left(e^T n(x, y) \right) L,$$

$a(x, y)$ surface albedo, L light source intensity. Superposition of arbitrary number of point sources at infinity is still in same 3D linear subspace, due to linear superposition of each contribution to image. Fisher images can eliminate within-class scatter.

Appearance manifold approach

For everyobject, (1) sample the set of viewing conditions (2) use these images as feature vectors (3) apply a PCA over all the images (4) keep the dominant PCs (5) sequence of views for one object represent a manifold in space of projections (6) what is the nearest manifold for a given view?


Object-pose manifold Appearance changes projected on PCs (1D pose changes). Sufficient characterization for recognition and pose estimation.

6.5 JPEG image compression

We don't resolve high frequencies too well... let's use this to compress images ...JPEG!

Concept Block-based discrete cosine transform (DCT)

JPEG Encoding and Decoding Encoding:



Decoding:

DCT A variant of discrete fourier transform: Real numbers, fast implementation. Block sizes: (1) small block: faster, correlation exists between neighboring pixels (2) better compression in smooth regions The first coefficient $B(0,0)$ is the DC component, the average intensity. The top-left coefficients represent low frequencies, the bottom right high frequencies.

symbol	prob.	code	binary fraction
Z	0.5	1	0.1
Y	0.25	01	0.01
X	0.125	001	0.001
W	0.125	000	0.000

The code words, if regarded as a binary fraction, are pointers to the particular interval being coded. In Huffman code, the code words point to the base of each interval. The average code length is $H = -\sum p(s) \log_2 p(s) \rightarrow$ optimal.

7 Scale-space representations

From an original signal $f(x)$ generate a parametric family of signals $f^l(x)$, where fine-scale information is successively suppressed.

7.1 Image pyramid

Level 0: 1×1 , Level 1: 2×2 , Level 2: 4×4 , Level $J - 1$: $N/2 \times N/2$, Level J (base): $N \times N$.

7.2 Applications

(1) Search for correspondence: look at coarse scales, then refine with finer scales (2) Edge tracking: a "good" edge at a fine scale has parents at a coarser scale (3) Control of detail and computational cost in matching: e.g. finding stripes; terribly important in texture representation

7.3 Pyramids

Gaussian pyramid Smooth with gaussians, because "gaussian²" = another gaussian. Progressively blurred and subsampled versions of the image. Adds scale invariance to fixed-size algorithms.

Laplacian Pyramid $i++i$ Shown the information added in gaussian pyramid at each spatial scale. Useful for noise reduction & coding.

Wavelet/QMF Bandpassed representation, complete, but with aliasing and some non-oriented subbands. Recursive application of a two-band filter bank to the lowpass band of the previous stage yields octave band splitting.

Steerable pyramid Shown components at each scale and orientation separately. Non-aliased subbands. Good for texture and feature analysis.

7.4 Haar Transform

Two major sub-operations: (1) Scaling captures info at different frequencies (2) Translation captures info at different locations Can be represented by filtering and downsampling. Relatively poor energy compaction.

8 Optical Flow

In Visual Computing, people seem like to use

$$I_{\blacklozenge} = \frac{\partial I}{\partial \blacklozenge}, \quad u = \frac{\partial x}{\partial t}, \quad v = \frac{\partial y}{\partial t}.$$

8.1 Applications

① tracking ② structure from motion ③ stabilization ④ compression ⑤ Mo-saicing

8.2 Brightness constancy

Definition of Optical Flow "Apparent motion of brightness patterns". Ideally, the optical flow is the projection of the three-dimensional velocity vectors on the image.

Caution required ① Uniform, rotating sphere $\mathcal{O}\mathcal{F} = 0$ ② No motion, but changing lighting $\mathcal{O}\mathcal{F} \neq 0$

8.3 Mathematical formulation

$I(x, y, t)$ = brightness at (x, y) at time t .

Brightness constancy assumption:

$$I\left(\frac{dx}{dt}\hat{x} + \frac{dy}{dt}\hat{y} + \delta t, t + \delta t\right) = I(x, y, t)$$

Optical flow constraint equation:

$$\frac{dI}{dt} = \frac{\partial I}{\partial x} \frac{dx}{dt} + \frac{\partial I}{\partial y} \frac{dy}{dt} + \frac{\partial I}{\partial t} = 0$$

8.4 The aperture problem

The motion of an edge seen through an aperture is (in some cases) inherently ambiguous. E.g. the edge is physically moving upwards, but the edge motion alone is consistent with many other possible motions, and in this case the edge e.g. appears to move diagonally.

8.5 Optical Flow meaning

Estimate of observed projected motion field. Not always well defined! Compare: ① Motion Field (or Scene Flow), projection of 3-D motion field ② Normal Flow: observed tangent motion ③ Optic Flow: Apparent motion of the brightness pattern: Apparent motion of the brightness pattern (hopefully equal to motion field) ④ Consider barber pole illusion

Planar motion Ideal motions of a plane, X, Y being the horizontal and vertical direction and Z normal to the image plane: ① translation in X ② translation in Z ③ rotation around Z ④ rotation around Y

8.6 Regularization: Horn & Schunk algorithm

The Horn-Schunck algorithm assumes smoothness in the flow over the whole image, thus, it tries to minimize distortions in flow and prefers solutions which show more smoothness. The flow

is formulated as a global energy functional which is the sought to be minimized. This function is given for two-dimensional image streams as:

$$E = \iint dx dy \left[\left(\frac{\partial I}{\partial x} \frac{dx}{dt} + \frac{\partial I}{\partial y} \frac{dy}{dt} + \frac{\partial I}{\partial t} \right)^2 + \alpha^2 (\|\nabla \hat{x}\| + \|\nabla \hat{y}\|)^2 \right]$$

The associated ELE are

$$\frac{\partial L}{\partial \hat{x}} - \frac{\partial}{\partial x} \frac{\partial L}{\partial \frac{\partial \hat{x}}{\partial x}} - \frac{\partial}{\partial y} \frac{\partial L}{\partial \frac{\partial \hat{x}}{\partial y}} = 0,$$

$$\frac{\partial L}{\partial \hat{y}} - \frac{\partial}{\partial x} \frac{\partial L}{\partial \frac{\partial \hat{y}}{\partial x}} - \frac{\partial}{\partial y} \frac{\partial L}{\partial \frac{\partial \hat{y}}{\partial y}} = 0.$$

this gives

$$\frac{\partial I}{\partial x} \left(\frac{\partial I}{\partial x} \hat{x} + \frac{\partial I}{\partial y} \hat{y} + \frac{\partial I}{\partial t} \right) - \alpha^2 \Delta \hat{x} = 0,$$

$$\frac{\partial I}{\partial y} \left(\frac{\partial I}{\partial x} \hat{x} + \frac{\partial I}{\partial y} \hat{y} + \frac{\partial I}{\partial t} \right) - \alpha^2 \Delta \hat{y} = 0.$$

$$\text{with } \Delta = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}.$$

Remarks

① Coupled PDE solved using iterative methods and finite differences

$$\hat{x} = \Delta \hat{x} - \lambda \left(\frac{\partial I}{\partial x} \hat{x} + \frac{\partial I}{\partial y} \hat{y} + I \right) \frac{\partial I}{\partial x},$$

$$\hat{y} = \Delta \hat{y} - \lambda \left(\frac{\partial I}{\partial x} \hat{x} + \frac{\partial I}{\partial y} \hat{y} + I \right) \frac{\partial I}{\partial y}.$$

② More than two frames allow a better estimation of I . ③ Information spreads from corner-type patterns. ④ Errors at boundaries ⑤ Example of *regularisation*: selection principle for the solution of illposed problems.

8.7 Lucas-Kanade: Integrate over a Patch

Assume a single velocity for all pixels within a image patch $I = I(x, y)$

$$E(\hat{x}, \hat{y}) = \sum_{x, y \in \Omega} \left(\frac{\partial I}{\partial x} \hat{x} + \frac{\partial I}{\partial y} \hat{y} + I \right)^2$$

$$\frac{\partial E}{\partial \hat{x}} = 2 \sum \frac{\partial I}{\partial x} \left(\frac{\partial I}{\partial x} \hat{x} + \frac{\partial I}{\partial y} \hat{y} + I \right),$$

$$\frac{\partial E}{\partial \hat{y}} = 2 \sum \frac{\partial I}{\partial y} \left(\frac{\partial I}{\partial x} \hat{x} + \frac{\partial I}{\partial y} \hat{y} + I \right).$$

9 Questions

- $I_{\text{comp}} = I_a I_d + (1 - I_a) I_b$
- MAP, Maximum a posteriori detector.
- graph cuts
- Solve MRFs with graph cuts
- impulse respons $t(-x, -y)$
- Canny nonmaxima suppression
- Entropy Coding (Huffman code)
- Aperture problem: normal flow